

Chapter 1

Characterization of

Distributed Systems

Outline

- What is a distributed system?
- Examples of distributed systems
- Resource sharing
- Challenges

What is a distributed system?

- At a high-level of abstraction, a distributed system is one that is composed of multiple connected computers where a user requests services *by name*.
- In distributed systems, components interact solely via *message passing*.


What is a distributed system?

- Based on such layout, we observe the following significant consequences:
 - Concurrency is inherent in distributed systems,
 - There is no global clock,
 - Distributed systems fail in new ways.
- We **build** distributed systems for one main reason: *sharing resources*, all kinds of resources!

Selected application domains and associated networked applications


<i>Finance and commerce</i>	eCommerce e.g. Amazon and eBay, PayPal, online banking and trading
<i>The information society</i>	Web information and search engines, ebooks, Wikipedia; social networking: Facebook and MySpace.
<i>Creative industries and entertainment</i>	online gaming, music and film in the home, user-generated content, e.g. YouTube, Flickr
<i>Healthcare</i>	health informatics, on online patient records, monitoring patients
<i>Education</i>	e-learning, virtual learning environments; distance learning
<i>Transport and logistics</i>	GPS in route finding systems, map services: Google Maps, Google Earth
<i>Science</i>	The Grid as an enabling technology for collaboration between scientists
<i>Environmental management</i>	sensor technology to monitor earthquakes, floods or tsunamis

Examples of distributed systems: The Internet



- Few issues to observe:
 - The set of services is **open-ended**,
 - Backbones are links with high transmission capacities realized through **different technologies**,
 - The ability of the Internet today to handle multimedia data is currently very limited due to the lack of facilities to **reserve network bandwidth** for individual streams of data.

Examples of distributed systems: An Intranet



Copyright © 2022 Prof. Amr El-Kadi

- This is a portion of the Internet with clearly defined **boundaries** in which **local security polices** can be enforced.
- What is a *firewall* and how is it implemented?

- The main issues related to the deployment of components in an intranet are:
 - Users need to **share data**, so a distributed file service is needed.
 - Firewalls may be too restrictive, sometimes **fine-grained security** mechanisms are needed.
- The cost of software deployment and maintenance is a crucial issue; **proper system architecture** should be employed.

Mobile and ubiquitous computing

- Increasingly becoming the norm.
- Examples include:
 - Laptop computers.
 - Handheld devices, including Personal Digital Assistants (PDAs), mobile phones, and digital cameras.
 - Wearable devices, such as smart watches.
 - Embedded devices, such as hi-fi systems, microwave ovens, and vending machines.


- *Mobile computing* (also called *nomadic computing*) is defined as the ability to perform computing while roaming.
- Users are able to utilize resources such as printers at conveniently nearby places as they move around. This is known as *location-aware computing*.

- *Ubiquitous computing* is possible through the utilization of many small and cheap computing devices that are present in the users' environments, for example smart cards.
- Computational behavior will be limited by their physical function.

- Ubiquitous and mobile computing overlap since mobile users expect to benefit from computers that are everywhere.
- Yet they are distinct
 - Ubiquitous computing would benefit a user while he/she remains in a single environment.
 - Similarly, a mobile computing user can still perform computing tasks even if it involves only more sophisticated devices such as laptops and printers.

- Both mobile and ubiquitous computing raise **significant system issues** including: discovery of resources, automatic reconfiguration of mobile devices, coping with limited connectivity while on the move, and security.

Portable and handheld devices in a distributed system



Resource sharing

- Resource sharing varies widely in scope and in how close users collaborate together.
- Examples of the two extremes include: *search engines*, and *Computer-Supported Cooperative Working (CSCW)*.
- Shared resources are accessed by *services* via the *operations* they export.


- Getting a service in a distributed system is achieved via a *remote invocation mechanism*.
- In the simple case, we have two active entities (*client* and *server*).
- Not all distributed systems are constructed entirely using the client-server model.

The World Wide Web

- It started as an environment providing a community of physicists working for the European Center for Nuclear Research (CERN) with a medium for exchanging documents via the Internet.
- Its key feature is the structure it provides in *hypertext*.
- The Web is an *open system* for two main reasons:
 1. It is based on standardized communication and documents' structures.
 2. It is open with respect to the type of resources that it can manipulate via the use of “helper” applications and “plug-ins”.

- It is based on **three major** standard components:
 - The HyperText Markup Language (HTML).
 - Uniform Resource Locators (URLs).
 - A client-server model, with HyperText Transfer Protocol (HTTP).

Cloud computing



Challenges

- The design and development of distributed systems is **not** quite simple.
- Most of the issues we discuss here as **challenges** have been **resolved**, still designers have to be highly aware of them and take them into account.

1. Heterogeneity

- It applies to the following:
 - Networks
 - Computer hardware
 - Operating systems
 - Programming languages
 - Applications' interfaces

Middleware

- What is it?
- Most middleware deal with differences in operating systems and hardware (e.g., mobile code)
- They also provide uniformity for application developers and end-users

2. Openness

- In general, what is openness of a computer system?
- In distributed systems, this primarily targets how easy it is to add new resources to the system and make them available to clients.
- Openness cannot be achieved unless the key interfaces are published.

3. Security

- This has **three components**: confidentiality, integrity, and availability.
- Several security challenges are resolved through the use of encryption techniques.
- Examples of security challenges that have not been fully met include *denial of service attacks*, and *security of mobile code*.

4. Scalability

- What makes a system **scalable**?
- The Internet is a good example of a scalable system

Growth of the Internet (computers and web servers)

<i>Date</i>	<i>Computers</i>	<i>Web servers</i>	<i>Percentage</i>
1993, July	1,776,000	130	0.008
1995, July	6,642,000	23,500	0.4
1997, July	19,540,000	1,203,096	6
1999, July	56,218,000	6,598,697	12
2001, July	125,888,197	31,299,592	25
2003, July	~200,000,000	42,298,371	21
2005, July	353,284,187	67,571,581	19

- Designing scalable distributed systems present the following challenges:
 - Controlling the cost of physical resources
 - Controlling the performance loss
 - Preventing software resources running out
 - Avoiding performance bottlenecks
- Ideally, **no change** should be needed to the system as a result of a scale increase (or decrease)

5. Failure handling

- Failures in distributed systems are not the same as in other systems, they are *partial*.
- The following techniques exist for handling failures:
 - **Detecting Failures**: Some failures could be detected, but the real challenge is to continue operation when failures happen and cannot be detected, yet they are suspected.
 - **Masking Failures**: When you detect failures, they can be masked or made less severe.
 - **Tolerating Failures**: Inform clients of failure presence and let them tolerate!
 - **Recovering from Failures**: Snapshots and rolling back.
 - **Redundancy**: Replicate resources. The challenge here is keeping replicas synchronized without major performance degradation.

6. Concurrency

- Proper inter-process communication and synchronization mechanisms have to be utilized to ensure data consistency.

7. Transparency

- This is ***the*** key issue in a distributed system. The system should appear to end-users and application developers as a **whole** (one system) rather than a collection of interconnected components.

- Eight forms of transparency were identified:
 - *Access transparency*. Identical operations should be used to access local and remote resources.
 - *Location transparency*. Resources should be accessed without knowledge of their physical location.
 - *Concurrency transparency*. Ensures consistency of resources in the presence of concurrent access.
 - *Replication transparency*. End-users and applications developers should not be aware of the existence of multiple replicas of the resources.

- *Failure transparency*. End-users and applications developers should continue their tasks even with the presence of faults in hardware or software components.
- *Mobility transparency*. Both resources and clients should be allowed to move within the system without affecting the operation of the system.
- *Performance transparency*. The system should be reconfigured dynamically to match load variances.
- *Scaling transparency*. Allows the system and applications to scale dynamically with no change to the system structure or algorithms.